

*The subsequent sections constitutes the appendix.*

## A Additional results: Classical ML models

For comprehensiveness, we trained three classical ML models—SVM, Logistic Regression (LR), and Random Forest (RF)—using simple n-gram/tf-idf features on our causal relation classification task for both the biomedical and general domain datasets. Please note that we trained the classical ML models using the full training examples instead of 16 examples as in our few-shot settings. The result is summarized in Table 1.

**Table 1.** Evaluation results for open-domain dataset: SEMEVAL-2010 Task 8

	SVM			LR			RF		
	P	R	F1	P	R	F1	P	R	F1
COMAGC	71.4	25.0	37.0	100	10.0	18.1	100	15.0	26.1
GENE	60.0	15.0	24.0	0.0	0.0	0.0	50.0	50.0	9.00
DDI	59.3	95.0	73.0	59.3	95.0	73.0	60.0	90.0	72.0
SEMEVAL	75.0	30.0	42.8	63.6	35.0	45.1	52.1	60.0	55.8

The classical ML models performed relatively well on the DDI datasets, achieving an F1 score of 73.0, which outperformed the ICL model using GPT-3.5 (F1=68.9). However, on the other datasets, the classical ML models performed poorly compared to Transformer/LM-based models (Table 2 and 3 in the main paper), even when trained with the full training data. We believe this discrepancy is due to the relatively small size of the other datasets compared to the DDI datasets.

The challenge on using classical ML approach is that it often requires complex feature engineering, while in deep learning-based approach such as Transformer-based model, the high-dimensional features are learned automatically by the "deep"-layered model, often resulting in better accuracy.

## B Baseline (GPT-3.5-turbo-instruct) hyperparameters

In this work, we used OpenAI<sup>1</sup> API with `gpt-3.5-turbo-instruct` engine. In general we assume only query access to the LLMs (i.e., no gradients, no log probabilities). Table 2 summarizes the hyperparameter values for the ICL model experiment with `gpt-3.5-turbo-instruct`.

<sup>1</sup> <https://platform.openai.com/>

**Table 2.** Hyperparameter values.

<i>hyperparameter</i>	value
temperature	0.7
max token	100 ~ 400
top p	1
frequency penalty	0
presence penalty	0

## C MLM hyperparameters settings

We used the `biomed_roberta_base`<sup>2</sup> for all biomedical datasets (COMAGC, GENEC, and DDI), and `roberta-base`<sup>3</sup> for the SEMEVAL dataset. We used both models from the `huggingface` [1] models library. All models are implemented in Python with `Pytorch`<sup>4</sup> and `Transfomer`<sup>5</sup> library. The random seed of 203 is set for all experiments. Table 3 summarizes the hyperparameter values for fine-tuning the model experiments.

**Table 3.** MLM hyperparameter values.

<i>hyperparameter</i>	COMAGC	GENEC	DDI	SEMEVAL
batch size	8	8	8	8
max length	256	256	256	256
optimizer	Adam	Adam	Adam	Adam
lr	3e-5	3e-5	3e-5	3e-5
gradient acc. steps	4	4	4	4
adam eps.	1e-06	1e-06	1e-06	1e-06
warmup proportion	0.06	0.06	0.06	0.06
weight decay	0.01	0.01	0.01	0.01
max grad norm	1	1	1	1
epoch	30	30	30	20

## D CLM/decoder-only hyperparameters settings

In all evaluation experiments, the `bloomz-560m`<sup>6</sup> model from the `huggingface` models library were used for all datasets. All models are implemented in Python

<sup>2</sup> [https://huggingface.co/allenai/biomed\\_roberta\\_base](https://huggingface.co/allenai/biomed_roberta_base)

<sup>3</sup> <https://huggingface.co/FacebookAI/roberta-base>

<sup>4</sup> <https://pytorch.org/>

<sup>5</sup> <https://huggingface.co/docs/transformers/en/index>

<sup>6</sup> <https://huggingface.co/bigscience/bloomz-560m>

with Pytorch<sup>7</sup> and Transformer<sup>8</sup> library. The random seed of 203 is set for all experiments. Table 4 summarizes the hyperparameter values for fine-tuning the bloomz-560m model experiments.

**Table 4.** bloomz-560m hyperparameter values.

<i>hyperparameter</i>	COMAGC	GENEC	DDI	SEMEVAL
batch size	4	4	4	4
max length	256	256	256	256
optimizer	AdamW	AdamW	AdamW	AdamW
lr	3e-5	3e-5	3e-5	3e-5
gradient acc. steps	4	4	4	4
warmup proportion	0.06	0.06	0.06	0.06
max grad norm	1	1	1	1
epoch	30	10	15	20

## E Seq2SeqLM hyperparameters settings

In all evaluation experiments, the t5-base<sup>9</sup> model from the huggingface models library were used for training all datasets. All models are implemented in Python with Pytorch<sup>10</sup> and Transformer<sup>11</sup> library. The random seed of 203 is set for all experiments. Table 5 summarizes the hyperparameter values for fine-tuning the t5-base model experiments.

**Table 5.** t5-base hyperparameter values.

<i>hyperparameter</i>	COMAGC	GENEC	DDI	SEMEVAL
batch size	8	4	8	8
max length	256	256	256	256
optimizer	AdamW	AdamW	AdamW	AdamW
lr	3e-3	3e-4	3e-4	3e-4
gradient acc. steps	4	4	4	4
warmup proportion	0.06	0.06	0.06	0.06
max grad norm	1	1	1	1
epoch	30	20	20	20

<sup>7</sup> <https://pytorch.org/>

<sup>8</sup> <https://huggingface.co/docs/transformers/en/index>

<sup>9</sup> <https://huggingface.co/google-t5/t5-base>

<sup>10</sup> <https://pytorch.org/>

<sup>11</sup> <https://huggingface.co/docs/transformers/en/index>

## F Querying the KGs

We access the **Hetionet** KG through its official public Neo4j API<sup>12</sup>, with `neo4j.v1` Python library. Neo4j<sup>13</sup> is a third-party graph database that supports the Cypher language for querying and visualizing a knowledge graph. Hetionet also provides a public Neo4j browser app<sup>14</sup>.

As for the **Wikidata**, we access the KG through its official public SPARQL endpoint<sup>15</sup>, with `SPARQLWrapper` Python library. We employ the official wikidata API (e.g., `wbsearchentities` and `wbgetentities` functions) for extracting the Wikidata IDs for all variable pairs.

On Hetionet, we query up to 4 hops for extracting the KGs structures. However, for Wikidata, we query up to one hops to extract the KG structures, constrained by its huge sizes. To train a robust models that generalize well given any KG structures, we opted to not optimize the content from the KG structures to be included in the prompt. For instance, when there are more than  $m$  metapaths for a pair, we randomly select  $m$  of them,  $m$  being a *hyperparameter* of the number of metapaths to be included as prompt.

In all experiments, we include up to the following: 4 neighbors nodes, 5 common neighbors nodes, and 1 metapath, to be included in the prompt. These values are selected based on hyperparameters setting giving the best results in our experiments.

## G Prompt hyperparameters and example

(to-be-added)

## References

1. Wolf, T., Debut, L., Sanh, V., Chaumond, J., Delangue, C., Moi, A., Cistac, P., Rault, T., Louf, R., Funtowicz, M., Davison, J., Shleifer, S., von Platen, P., Ma, C., Jernite, Y., Plu, J., Xu, C., Le Scao, T., Gugger, S., Drame, M., Lhoest, Q., Rush, A.: Transformers: State-of-the-art natural language processing. In: Liu, Q., Schlangen, D. (eds.) Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations. pp. 38–45. Association for Computational Linguistics, Online (Oct 2020). <https://doi.org/10.18653/v1/2020.emnlp-demos.6>, <https://aclanthology.org/2020.emnlp-demos.6>

<sup>12</sup> <bolt://neo4j.het.io>

<sup>13</sup> <https://neo4j.com/>

<sup>14</sup> <https://neo4j.het.io/browser/>

<sup>15</sup> <https://query.wikidata.org/sparql>